

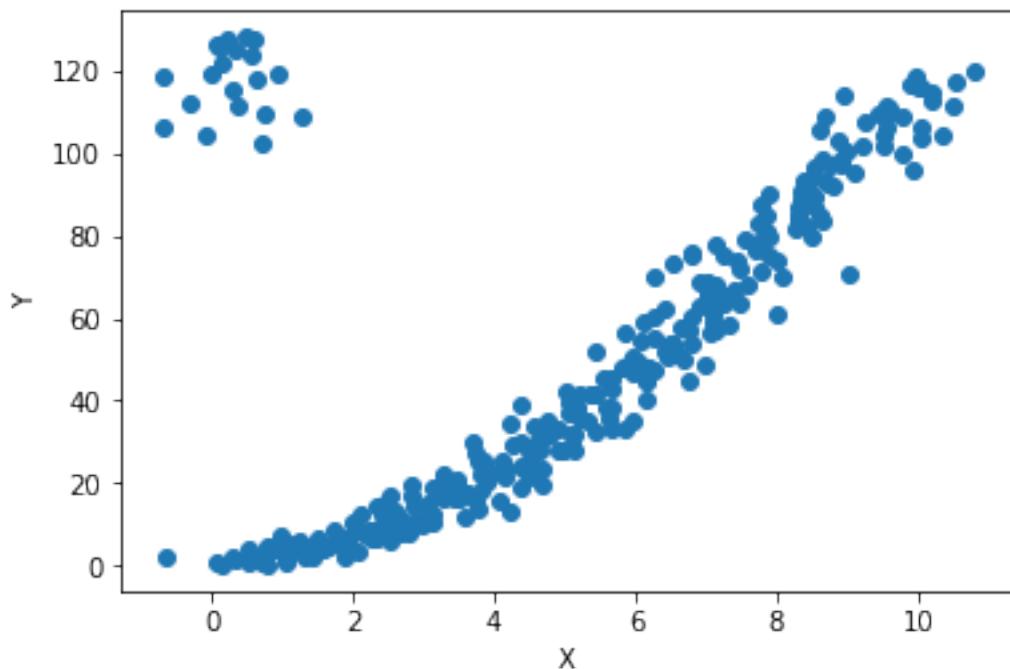
20180213_COGS118a_Hw4

February 14, 2018

1 Parabola Estimation

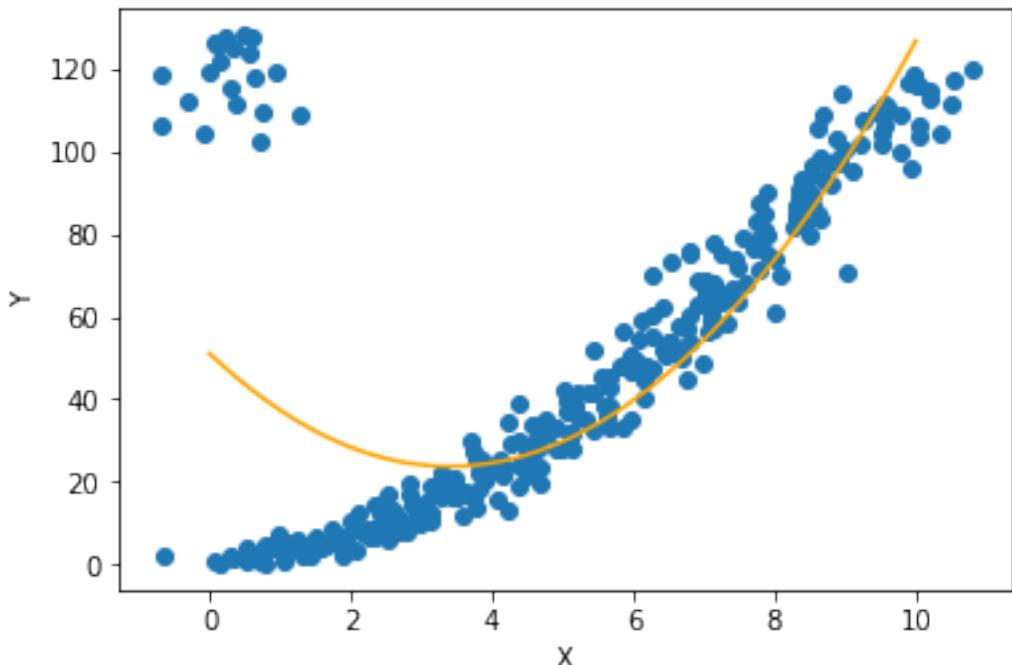
```
In [1]: import numpy as np
        import matplotlib.pyplot as plt
        X_and_Y = np.load('./hw4-q1-parabola.npy')
        X = X_and_Y[:, 0]
        Y = X_and_Y[:, 1]

        plt.scatter(X, Y)
        plt.xlabel('X')
        plt.ylabel('Y')
        plt.show()
```



```
In [12]: X1 = np.matrix(np.hstack((np.ones((len(X),1)),
X.reshape(-1,1))))  
  
X2 = np.matrix(np.hstack((X1, (X**2).reshape(-1,1))))  
W = X2.T.dot(X2).I.dot(X2.T).dot(Y)  
L2w0, L2w1, L2w2 = np.array(W).reshape(-1)  
print('Y = {:.2f} + {:.2f}*X + {:.2f}*X2'.format(w0, w1, w2))  
  
Y = 50.75 + -15.94*X + 2.35*X2
```

```
In [13]: X_line = np.linspace(0,10,300)  
Y_line = L2w0 + L2w1 * X_line + L2w2 * (X_line**2)  
plt.scatter(X, Y)  
plt.plot(X_line, Y_line, color='orange')  
plt.xlabel('X')  
plt.ylabel('Y')  
plt.show()
```



```
In [22]: # g'(W)  
def g_prime_W(X, Y, W):  
    return (np.sign(X.dot(W) - Y).T.dot(X)).T  
  
W = np.matrix(np.zeros((3,1)))  
Y = Y.reshape(-1, 1)
```

By the chain rule,

$$\frac{\partial |f(w)|}{\partial w} = \text{sign}(f(w)) \cdot \frac{\partial f(w)}{\partial w}$$

Hence,

$$\begin{aligned}\frac{\partial g(w)}{\partial w} &= \sum_{i=1}^n \text{sign}(x_i w_i - y_i) \cdot x_i \\ &= \begin{bmatrix} \text{sign}(x_1 w_1 - y_1) \cdot 1 + \dots + \text{sign}(x_n w_n - y_n) \cdot 1 \\ \text{sign}(x_1 w_1 - y_1) \cdot x_1 + \dots + \text{sign}(x_n w_n - y_n) \cdot x_n \\ \text{sign}(x_1 w_1 - y_1) x_1^2 + \dots + \text{sign}(x_n w_n - y_n) x_n^2 \end{bmatrix}\end{aligned}$$

By the identity that $(a_1 \dots a_n) \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix} = \sum_{i=1}^n a_i b_i$,
the above can be written as,

$$\text{sign}(Xw - Y)^T X$$

Because the shape of the gradient must match
the shape of the variable we take the gradient
w.r.t., we take the transpose of the above,

$$(\text{sign}(Xw - Y)^T X)^T$$

1b

```
# We will keep track of training loss over iterations
iterations = [0]
g_W = [np.absolute(X2.dot(W) - Y)]
for i in range(300000):
    grad = g_prime_W(X2, Y, W)
    W_new = W - 0.000001 * grad
    iterations.append(i+1)
    g_W.append((X2.dot(W_new) - Y).T.dot(X2.dot(W_new) - Y))
    if np.linalg.norm(W_new - W, ord = 1) < 0.00001:
        print("gradient descent terminated after " + str(i) + " iterations")
        break
    W = W_new
L1w0, L1w1, L1w2 = np.array(W).reshape(-1)
print('Y = {:.2f} + {:.2f}*X1 + {:.2f}*X2'.format(L1w0, L1w1, L1w2))
```

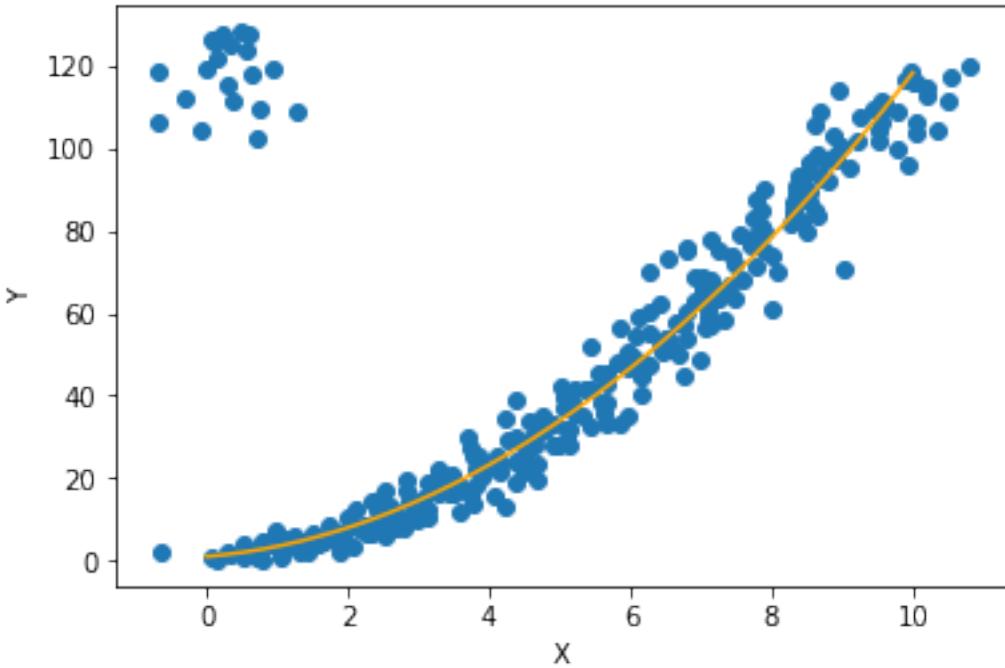
gradient descent terminated after 41383 iterations
 $Y = 1.07 + 1.43*X1 + 1.03*X2$

```
In [23]: X_line = np.linspace(0, 10, 300)
Y_line = L1w0 + L1w1 * X_line + L1w2 * (X_line**2)
plt.scatter(X, Y)
plt.plot(X_line, Y_line, color='orange')
plt.xlabel('X')
```

```

plt.ylabel('Y')
plt.show()

```



```

In [24]: # g'(W)
def g_prime_Walpha(X, Y, W, alpha):
    return alpha * (X.T.dot(2 * (X.dot(W) - Y))) + (1-alpha) * ((np.sign(X.dot(W) - Y)

alpha_dict = {}
for alpha in [0.3, 0.5, 0.7]:
    W = np.matrix(np.zeros((3,1)))
    Y = Y.reshape(-1, 1)
    # We will keep track of training loss over iterations
    iterations = [0]
    g_W = [alpha * ((X2.dot(W) - Y).T.dot(X2.dot(W) - Y)) + (1 - alpha) * np.absolute(X2
    for i in range(300000):
        grad = g_prime_Walpha(X2, Y, W, alpha)
        W_new = W - 0.000001 * grad
        iterations.append(i+1)
        g_W.append((X2.dot(W_new) - Y).T.dot(X2.dot(W_new) - Y))
        if np.linalg.norm(W_new - W, ord = 1) < 0.00001:
            print("gradient descent terminated after " + str(i) + " iterations")
            break
        W = W_new
    alpha_dict[alpha] = np.array(W).reshape(-1)
print(alpha_dict)

```

```

gradient descent terminated after 236048 iterations
{0.3: array([ 49.80810639, -15.58382493,    2.32721435])}
gradient descent terminated after 157048 iterations
{0.3: array([ 49.80810639, -15.58382493,    2.32721435]), 0.5: array([ 50.46529152, -15.830045 ,   2.32721435])}
gradient descent terminated after 119444 iterations
{0.3: array([ 49.80810639, -15.58382493,    2.32721435]), 0.5: array([ 50.46529152, -15.830045 ,   2.32721435])}

```

```

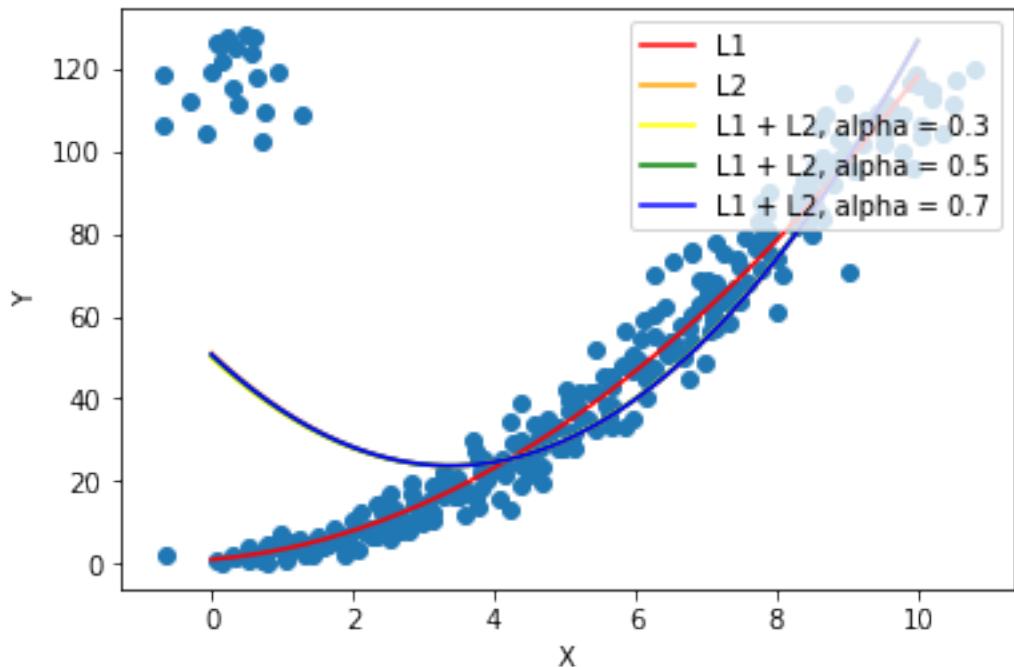
In [25]: X_line = np.linspace(0,10,300)
L1Y_line = L1w0 + L1w1 * X_line + L1w2 * (X_line**2)
L2Y_line = L2w0 + L2w1 * X_line + L2w2 * (X_line**2)
alpha3Y_line = alpha_dict[.3][0] + alpha_dict[.3][1] * X_line + alpha_dict[.3][2] * (X_line**2)
alpha5Y_line = alpha_dict[.5][0] + alpha_dict[.5][1] * X_line + alpha_dict[.5][2] * (X_line**2)
alpha7Y_line = alpha_dict[.7][0] + alpha_dict[.7][1] * X_line + alpha_dict[.7][2] * (X_line**2)

plt.scatter(X, Y)

plt.plot(X_line, L1Y_line, color='red', label='L1')
plt.plot(X_line, L2Y_line, color='orange', label='L2')
plt.plot(X_line, alpha3Y_line, color='yellow', label='L1 + L2, alpha = 0.3')
plt.plot(X_line, alpha5Y_line, color='green', label='L1 + L2, alpha = 0.5')
plt.plot(X_line, alpha7Y_line, color='blue', label='L1 + L2, alpha = 0.7')

plt.xlabel('X')
plt.ylabel('Y')
plt.legend(loc = 'upper right')
plt.show()

```



1. The reason that the L1 loss function is a better fit to the data than the L2 loss function is because the L2 loss function is much more sensitive to outliers, which can be seen in the upper left quadrant of the graph.
2. The L2 curve appears to have a sharper curve than the L1 curve because of the outliers.
3. The L2 curve is similar to the L1 + L2 curves because the square term in the L1 + L2 equations dominate the behavior of the graph, hence they will look similar to the L2 curve.

hw4-q2-logistic

February 14, 2018

0.1 Q2 Logistic Regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
import math
%config InlineBackend.figure_format = 'retina'
%matplotlib inline
from sklearn.utils import shuffle
import scipy.io as sio
plt.rcParams['figure.figsize'] = 8,8
```

0.1.1 Original Data

```
In [2]: X_and_Y_train = np.load('./hw4-q2-logistic-train.npy')
X_train = X_and_Y_train[:, :2]      # Shape: (70,2)
X_train = np.matrix(np.hstack((np.ones((len(X_train), 1)), X_train))) # Shape: (70, 3)
Y_train = X_and_Y_train[:, 2]       # Shape: (70,)

In [3]: X_and_Y_test = np.load('./hw4-q2-logistic-test.npy')
X_test = X_and_Y_test[:, :2]       # Shape: (30,2)
X_test = np.matrix(np.hstack((np.ones((len(X_test), 1)), X_test))) # Shape: (30, 3)
Y_test = X_and_Y_test[:, 2]        # Shape: (70,)

In [4]: mpl.style.use('seaborn')

fig = plt.figure()
plt.scatter([X_train[Y_train==0, 1]], [X_train[Y_train==0, 2]], marker='x', color='b', alpha=0.5)
plt.scatter([X_train[Y_train==1, 1]], [X_train[Y_train==1, 2]], marker='o', color='r', alpha=0.5)
plt.xlabel('X1')
plt.ylabel('X2')
plt.legend(loc='upper right', fontsize=10)
plt.title('Training data')
plt.show()
#fig.savefig('scatter_1.png', format='png', dpi=400)
```

$$\mathcal{L}(\theta) = - \sum_i \left[y^{(i)} \ln h(x^{(i)}; \theta) + (1-y^{(i)}) \ln (1-h(x^{(i)}; \theta)) \right]$$

Let $z = f(x^{(i)}; \theta) = \sum_{k=0}^K \theta_k x_k$

By the Chain Rule,

$$\frac{\partial \mathcal{L}(\theta)}{\partial \theta} = - \sum \left[y^{(i)} \frac{1}{h(z)} + (1-y^{(i)}) \frac{1}{1-h(z)} \right] \frac{\partial}{\partial \theta} h(z)$$

Again by the Chain Rule,

$$\begin{aligned} &= - \sum \left[y^{(i)} \frac{1}{h(z)} + (1-y^{(i)}) \frac{1}{1-h(z)} \right] h(z) \left(\frac{\partial}{\partial \theta} z \right) \\ &= - \sum \left[y^{(i)} (1-h(z)) + (1-y^{(i)}) h(z) \right] x^{(i)} \\ &= - \sum \left[y^{(i)} - h(z) \right] x^{(i)} \\ \therefore &= - \sum \left[y^{(i)} - h(x^{(i)}; \theta) \right] x^{(i)} \end{aligned}$$



0.1.2 Gradient Descent

```
In [5]: def sigmoid(z):
    return 1.0/(1.0+np.exp(-z))
```

```
In [6]: # gradient of loss function L(theta)
def L_prime_theta(X, Y, theta):
    Y = Y.reshape(-1, 1)
    retVal = -((Y - sigmoid(X.dot(theta))).T.dot(X)).T
    return retVal
```

```
In [7]: def L_theta(X, Y, theta):
    return -(Y.dot(np.log(sigmoid(X.dot(theta)))) + (1 - Y).dot(np.log(1 - sigmoid(X.dot(theta)))))
```

```
In [8]: learning_rate = 0.001
n_iter = 10000
theta = np.zeros((X_train.shape[1], 1))
# We will keep track of training loss over iterations
iterations = [0]
L_theta_list = [L_theta(X_train, Y_train, theta)]
for i in range(n_iter):
    gradient = L_prime_theta(X_train, Y_train, theta)
```

```

theta_new = theta - learning_rate * gradient
iterations.append(i+1)
L_theta_list.append(L_theta(X_train, Y_train, theta_new))

if np.linalg.norm(theta_new - theta, ord = 1) < 0.001:
    print("gradient descent has converged after " + str(i) + " iterations")
    break
theta = theta_new

print ("theta vector: \n" + str(theta))

gradient descent has converged after 6365 iterations
theta vector:
[[-11.48322099]
 [ 0.97466299]
 [ 0.88907048]]

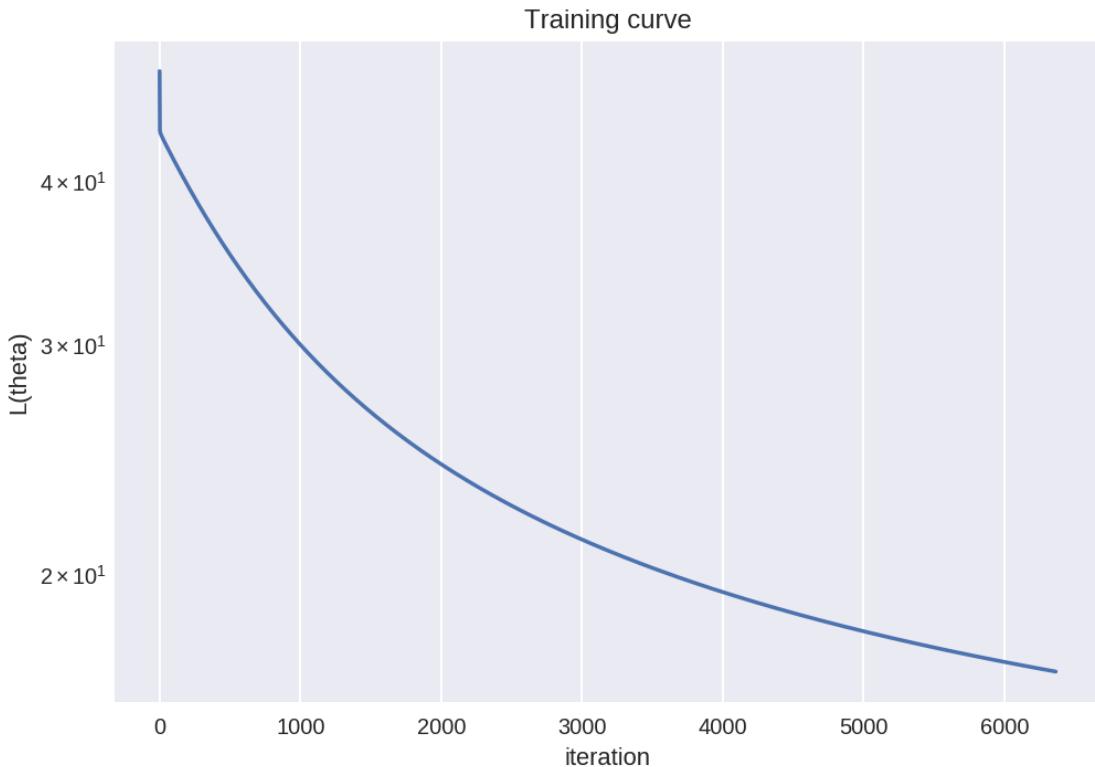
```

Equation of decision boundary corresponding to optimal θ^* :

$$y = \begin{cases} 1 & -11.4832 + 0.9747x_1 + 0.8891x_2 \geq 0 \\ 0 & -11.4832 + 0.9747x_1 + 0.8891x_2 < 0 \end{cases}$$

0.1.3 Training curve

```
In [9]: plt.title('Training curve')
plt.xlabel('iteration')
plt.ylabel('L(theta)')
plt.semilogy(iterations, np.array(L_theta_list).reshape(-1, 1))
plt.show()
```



0.1.4 Results on Training data

```
In [10]: prediction = sigmoid(np.dot(X_train, theta)) >= 0.5
testing_accuracy = np.sum(prediction == Y_train.reshape(-1, 1))*1.0/X_train.shape[0]
print(prediction.shape, Y_test.shape)
print ("training accuracy: " + str(testing_accuracy))

x = np.arange(np.min(X_train[:,1])-1,np.max(X_train[:,1])+1,1.0)
y = (-theta[0][0]-theta[1][0]*x)/theta[2][0]
plt.scatter([X_train[Y_train==0, 1]], [X_train[Y_train==0, 2]], marker='x', color='b',
plt.scatter([X_train[Y_train==1, 1]], [X_train[Y_train==1, 2]], marker='o', color='r',

plt.xlabel('X1')
plt.ylabel('X2')
plt.plot(x,y.T, 'dodgerblue', label='decision boundary')
plt.title('Training data and decision boundary')

plt.legend(loc='upper right', fontsize=10)

(70, 1) (30,)
training accuracy: 0.914285714286
```

```
Out[10]: <matplotlib.legend.Legend at 0x7f512a5ed9e8>
```



0.1.5 Results on Testing data

```
In [11]: prediction = sigmoid(np.dot(X_test, theta)) >= 0.5
testing_accuracy = np.sum(prediction == Y_test.reshape(-1, 1))*1.0/X_test.shape[0]
print ("testing accuracy: " + str(testing_accuracy))

x = np.arange(np.min(X_train[:,1])-1,np.max(X_train[:,1])+1,1.0)
y = (-theta[0][0]-theta[1][0]*x)/theta[2][0]
plt.scatter([X_test[Y_test==0, 1]], [X_test[Y_test==0, 2]], marker='x', color='b', alpha=0.5)
plt.scatter([X_test[Y_test==1, 1]], [X_test[Y_test==1, 2]], marker='o', color='r', alpha=0.5)

plt.xlabel('X1')
plt.ylabel('X2')
plt.plot(x,y.T, 'dodgerblue', label='decision boundary')
plt.title('Testing data and decision boundary')

plt.legend(loc='upper right', fontsize=10)

testing accuracy: 0.833333333333
```

Out[11]: <matplotlib.legend.Legend at 0x7f5128d59898>



hw4-q3-lda

February 14, 2018

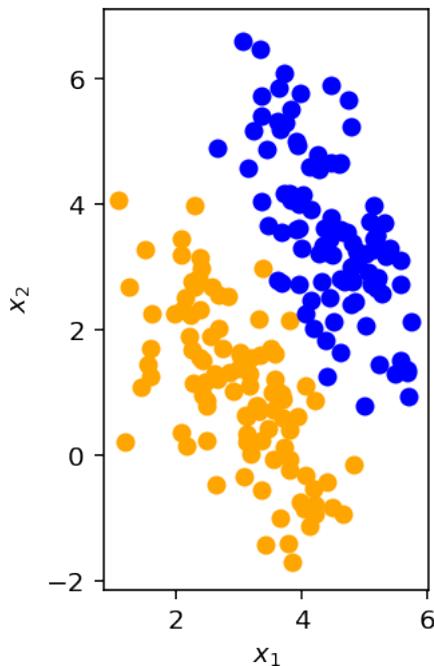
```
In [2]: import numpy as np
        import matplotlib.pyplot as plt
        %config InlineBackend.figure_format = 'retina'

In [3]: # Load the data and visualize.
Xs = np.load('hw4-q3-lda.npy')

X_0 = np.matrix(Xs[:, 0:2]).T # Shape: (2, 100).
X_1 = np.matrix(Xs[:, 2:4]).T # Shape: (2, 100).

print(X_0.shape, X_1.shape)
plt.scatter(X_0[0].tolist(), X_0[1].tolist(), color='orange')
plt.scatter(X_1[0].tolist(), X_1[1].tolist(), color='blue')
plt.axis('scaled')
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.show()

(2, 100) (2, 100)
```



```
In [12]: # (a) Compute mean of each class.
mu_0 = np.matrix([X_0[0].mean(), X_0[1].mean()]).T # Shape: (2, 1).
mu_1 = np.matrix([X_1[0].mean(), X_1[1].mean()]).T # Shape: (2, 1).

print(mu_0.shape, mu_1.shape)
print('mu_0=\n{} ,\nmu_1=\n{}'.format(mu_0, mu_1))

(2, 1) (2, 1)
mu_0=
[[ 2.98351552]
 [ 1.06453902]],
mu_1=
[[ 4.46952033]
 [ 3.52885988]]
```



```
In [14]: # (b) Compute the covariance matrix for each class, Sigma_0 and Sigma_1.
Sigma_0 = np.cov(X_0) # Shape: (2, 2).
Sigma_1 = np.cov(X_1) # Shape: (2, 2).

print(Sigma_0.shape, Sigma_1.shape)
print('Sigma_0=\n{} ,\nSigma_1=\n{}'.format(Sigma_0, Sigma_1))

(2, 2) (2, 2)
Sigma_0=
```

```
[[ 0.70628859 -0.6905174 ]
 [-0.6905174   1.61474336]],
Sigma_1=
[[ 0.48981843 -0.57477756]
 [-0.57477756  1.67666167]]
```

In [16]: `from numpy import linalg as LA`

```
# (c) Find the optimal w_star and w_tilde_star with unit length.
w_star      = LA.inv(Sigma_0 + Sigma_1).dot(mu_0 - mu_1) # Shape: (2, 1).
w_tilde_star = w_star/LA.norm(w_star, 2) # Shape: (2, 1).

print(w_star.shape, w_tilde_star.shape)
print('w_star=\n{}\n,w_tilde_star=\n{}'.format(w_star, w_tilde_star))

(2, 1) (2, 1)
w_star=
[[-3.42871346]
 [-2.06679356]],
w_tilde_star=
[[-0.85643702]
 [-0.51625152]]
```

In [19]: # (d) Compute the projection and plot the figure.

```
Xproj_0 = w_tilde_star.T.dot(X_0).T.dot(w_tilde_star.T).T # Shape: (2, 100).
Xproj_1 = w_tilde_star.T.dot(X_1).T.dot(w_tilde_star.T).T # Shape: (2, 100).

print(Xproj_0.shape, Xproj_1.shape)
plt.scatter(X_0[0].tolist(), X_0[1].tolist(), color='orange')
plt.scatter(X_1[0].tolist(), X_1[1].tolist(), color='blue')
plt.scatter(Xproj_0[0].tolist(), Xproj_0[1].tolist(), color='yellow')
plt.scatter(Xproj_1[0].tolist(), Xproj_1[1].tolist(), color='green')
plt.axis('scaled')
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.show()

(2, 100) (2, 100)
```

